

Capstone Project

Image classifier for the SVHN dataset

Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
import tensorflow.keras.layers as layers
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping

In [2]: # This is needed to run on Windows 10 laptop with Cuda
print(tf.__version__)

config = tf.compat.v1.ConfigProto()
config.gpu_options.allow_growth = True
session = tf.InteractiveSession(config=config)

2.1.0
```

SVHN Overview image

For the capstone project you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

Y. N. Netzer, D. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning".
NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

```
In [3]: # Run this cell to load the dataset
train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. Hint: retain the channel dimension, which will now have size 1.
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
In [4]: # Task:
# Extract the training and testing images and labels separately from the train and test dictionaries
# loaded for you.
import numpy as np
X_train = train['X']
y_train = train['y']
X_test = test['X']
y_test = test['y']
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)

(32, 32, 3, 73257)
(32, 32, 1, 73257)
(32, 32, 3, 26032)
(32, 32, 1, 26032)
```

```
In [5]: from sklearn.preprocessing import OneHotEncoder

enc = OneHotEncoder().fit(y_train)
y_train = enc.transform(y_train).toarray()
y_test = enc.transform(y_test).toarray()
```

```
In [6]: print(y_train.shape)

(73257, 10)
```

```
In [7]: # Task:
# Convert the training and test images to grayscale by taking the average across all colour channels
# for each pixel. Hint: retain the channel dimension, which will now have size 1
```

```
In [8]: X_train_avg = np.swapaxes(X_train, 3, 0)
print(X_train_avg.shape)
X_train_avg = np.swapaxes(X_train_avg, 3, 2)
X_train_avg = np.swapaxes(X_train_avg, 1, 2)
print(X_train_avg.shape)

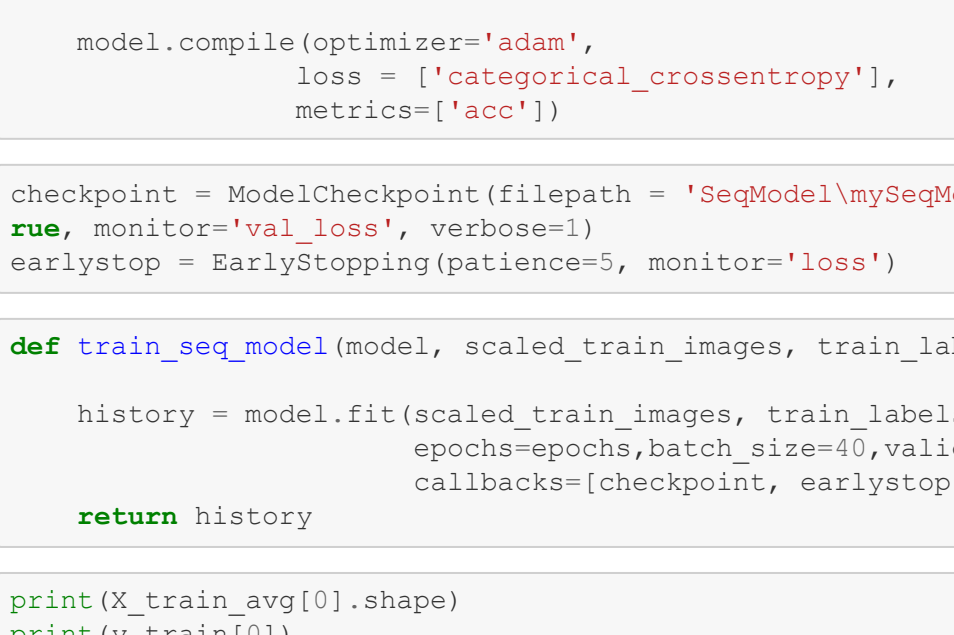
(73257, 32, 32, 3)
(73257, 32, 32, 3)
```

```
In [9]: X_train_avg = np.mean(X_train_avg,axis=3)
X_train_avg = X_train_avg[...,:np.newaxis]
print(X_train_avg.shape)
plt.imshow(X_train_avg[1].shape)

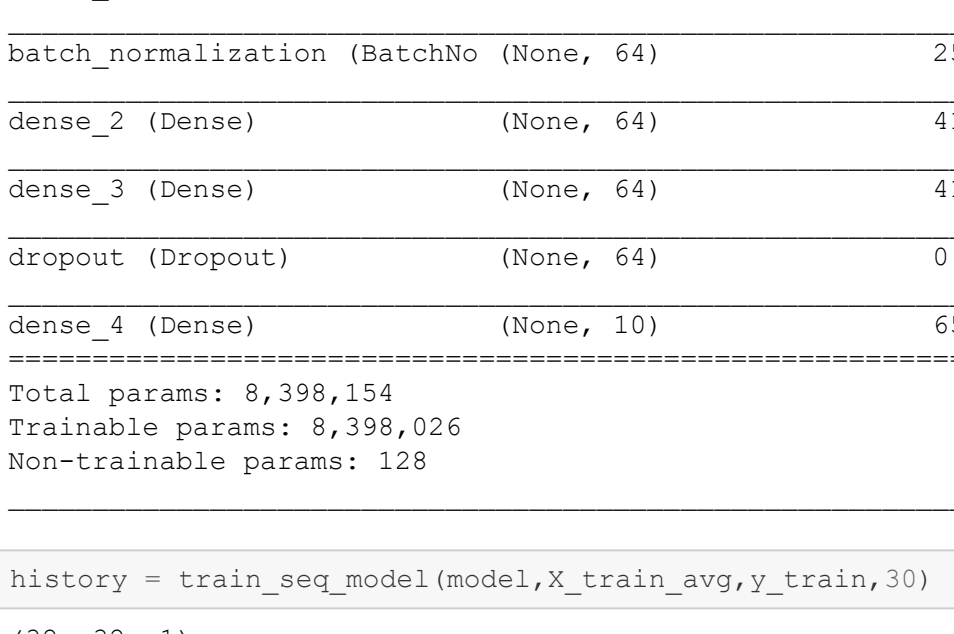
(73257, 32, 32, 1)
```

```
In [10]: # Task:
# Select a random sample of images and corresponding labels from the dataset (at least 10), and display
# them in a figure.
import numpy as np
import matplotlib.pyplot as plt
import random

w=10
h=10
fig=plt.figure(figsize=(8,4))
columns = 5
rows = 2
shp=X_train_avg.shape[0]
for i in range(1, columns*rows+1):
    ind = random.randint(0,shp)
    img = X_train_avg[ind,:].squeeze(axis=2)
    fig.add_subplot(rows,columns, i)
    plt.imshow(img)
```



```
In [11]: # Task:
# Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display
# them in a figure.
print(X_train_avg[1].squeeze(axis=2).shape)
w=10
h=10
fig=plt.figure(figsize=(8,4))
columns = 5
rows = 2
shp=X_train_avg.shape[0]
for i in range(1, columns*rows+1):
    ind = random.randint(0,shp)
    img = X_train_avg[ind,:].squeeze(axis=2)
    fig.add_subplot(rows,columns, i)
    plt.imshow(img, cmap='gray', vmin=0, vmax=255, label=y_train[ind])
    plt.show()
```



2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- Use accuracy and build the model yourself. For 4 weeks to experiment with different MLP architectures. Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.
- Print out the model summary (using the summary() method).
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.

- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).

- Plot the learning curves for loss and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [12]: def get_seq_model(input_shape):
model = Sequential([
    Dense(128, input_shape=input_shape, activation='relu'),
    Flatten(),
    Dense(64, activation='relu'),
    BatchNormalization(),
    Dense(64, activation='relu'),
    Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    Dense(10, activation='softmax'),
])
return model
```

```
In [13]: def compile_seq_model(model):
model.compile(optimizer='adam',
              loss=['categorical_crossentropy'],
              metrics=['acc'])
```

```
In [14]: checkpoint = ModelCheckpoint(filepath = 'SeqModel\mySeqModel', save_best_only=True, save_weights_only=True,
monitor='val_loss', verbose=1)
earlystop = EarlyStopping(patience=5, monitor='loss')
```

```
In [15]: def train_seq_model(model, scaled_train_images, train_labels, epochs):
history = model.fit(scaled_train_images, train_labels,
                    epochs=epochs, batch_size=32, validation_split=0.15, verbose=1,
                    callbacks=[checkpoint, earlystop])
return history
```

```
In [16]: print(X_train_avg[0].shape)
print(y_train[0])
model = get_seq_model(X_train_avg[0].shape)
compile_seq_model(model)
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 128)	16384
flatten (Flatten)	(None, 1024)	0
dense_2 (Dense)	(None, 64)	65536
batch_normalization (Batch Normalization)	(None, 64)	256
dense_3 (Dense)	(None, 64)	4160
dense_4 (Dense)	(None, 64)	4160
dropout (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 10)	650

Total params: 8,398,154
Trainable params: 8,398,026
Non-trainable params: 128

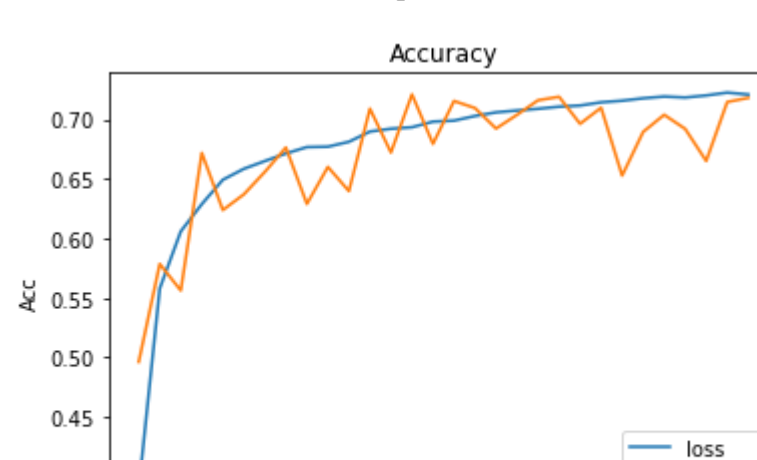
```
In [17]: history = train_seq_model(model,X_train_avg,y_train,30)

(32, 32, 1)
[[1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]]
Train on 62268 samples, validate on 10989 samples
Epoch 1/30
62200/62268 [=====] - ETA: 0s - loss: 1.7648 - acc: 0.3899
Epoch 00001: val_loss improved from inf to 1.52954, saving model to SeqModel\mySeqModel
62268/62268 [=====] - 458 765us/sample - loss: 1.7644 - acc: 0.3900 - val_loss
ss: 1.5295 - val_acc: 0.4964
Epoch 2/30
62200/62268 [=====] - ETA: 0s - loss: 1.3678 - acc: 0.5580
Epoch 00002: val_loss improved from 1.52954 to 1.25956, saving model to SeqModel\mySeqModel
62268/62268 [=====] - 468 731us/sample - loss: 1.3677 - acc: 0.5580 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 3/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00003: val_loss did not improve from 1.25956
62268/62268 [=====] - 458 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 4/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00004: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 5/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00005: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 6/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00006: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 7/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00007: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 8/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00008: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 9/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00009: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 10/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00010: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 11/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00011: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 12/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00012: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 13/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00013: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 14/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00014: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 15/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00015: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 16/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00016: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 17/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00017: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 18/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00018: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 19/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00019: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 20/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00020: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 21/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00021: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 22/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00022: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 23/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00023: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 24/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00024: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 25/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00025: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 26/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00026: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 27/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00027: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 28/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00028: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 29/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00029: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
Epoch 30/30
62200/62268 [=====] - ETA: 0s - loss: 1.2472 - acc: 0.6059
Epoch 00030: val_loss did not improve from 1.25956
62268/62268 [=====] - 468 731us/sample - loss: 1.2471 - acc: 0.6059 - val_loss
ss: 1.2596 - val_acc: 0.5786
```

```
In [18]: model.save('seqmodel\seqmodel.h5')
```

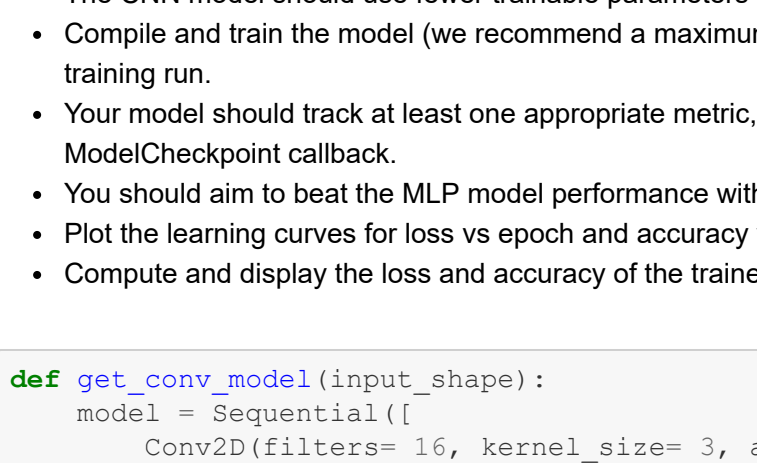
```
In [19]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['loss', 'val_loss'], loc='upper right')
plt.title('Loss')
```

```
Out [19]: Text(0.5, 1.0, 'Loss')
```



```
In [22]: plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.xlabel('Epochs')
plt.ylabel('Acc')
plt.legend(['loss', 'val_acc'], loc='lower right')
plt.title('Accuracy')
```

```
Out [22]: Text(0.5, 1.0, 'Accuracy')
```



3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, and Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.

- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [23]: def get_conv_model(input_shape):
model = Sequential([
    Conv2D(filters=32, kernel_size=(3,3), activation='relu', input_shape=input_shape),
    MaxPooling2D(pool_size=(3,3), strides=1),
    Conv2D(filters=32, kernel_size=(3,3), padding='valid', strides=1, activation='relu'),
    MaxPooling2D(pool_size=(1,1), strides=3),
    BatchNormalization(),
    Conv2D(filters=32, kernel_size=(3,3), padding='valid', strides=1, activation='relu'),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    Dense(10, activation='softmax'),
])
return model
```

```
In [24]: def compile_conv_model(model):
model.compile(optimizer='adam',
              loss=['categorical_crossentropy'],
              metrics=['accuracy'])
```

```
In [25]: from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
def get_early_stopping():
"""
This function should return an EarlyStopping callback that stops training when
the validation (testing) accuracy has not improved in the last 7 epochs.
HINT: use the EarlyStopping callback with the correct 'patience' and 'monitor'
"""
early_stopping = tf.keras.callbacks.EarlyStopping(patience=7, monitor='val_accuracy', mode='max')
return early_stoppingcallback
```

```
def get_checkpoint_best_only():
"""
This function should return a ModelCheckpoint object that:
- saves only the weights that generate the highest validation (testing) accuracy
- saves into a directory called 'checkpoints' (using the current working directory)
- generates a file called 'checkpoints_best_only/checkpoint'
"""
checkpoint_path = 'checkpoints_best_only/checkpoint'
checkpoint = ModelCheckpoint(filepath=checkpoint_path,
                             save_weights_only=False,
                             save_freq='epoch',
                             monitor='val_accuracy',
                             save_best_only=True,
                             verbose=1)
return checkpoint
```

```
In [26]: checkpoint_best_only = get_checkpoint_best_only()
early_stopping = get_early_stopping()
```

```
In [27]: def train_conv_model(model, scaled_train_images, train_labels, epochs):
history = model.fit(scaled_train_images, train_labels,
                    epochs=epochs, batch_size=32, validation_split=0.15, verbose=1,
                    callbacks=[checkpoint_best_only, early_stopping])
return history
```

```
In [28]: model2 = get_conv_model(X_train_avg[0].shape)
compile_conv_model(model2)
history = train_conv_model(model2,X_train_avg,y_train,30)
```

Train on 62268 samples, validate on 10989 samples
Epoch 1/30
62080/62268 [=====] - ETA: 0s - loss: 1.4067 - accuracy: 0.5260
Epoch 00001: val_accuracy improved from -inf to 0.69451, saving model to checkpoints_best_only/checkpoint

WARNING:tensorflow:From C:\Users\omkar\karve\.conda\envs\ml\lib\site-packages\tensorflow\python\ops\resource_variable_ops.py:1786: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

If using Keras pass 'constraint arguments to layers.
INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 208 328us/sample - loss: 1.4055 - accuracy: 0.5265 - v

al_loss: 1.0636 - val_accuracy: 0.6945
Epoch 2/30
62200/62268 [=====] - ETA: 0s - loss: 0.7886 - accuracy: 0.7603
Epoch 00002: val_accuracy improved from 0.69451 to 0.84029, saving model to checkpoints_best_only/checkpoint

INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 188 285us/sample - loss: 0.7888 - accuracy: 0.7603 - v

al_loss: 0.5371 - val_accuracy: 0.8403
Epoch 3/30
62200/62268 [=====] - ETA: 0s - loss: 0.6725 - accuracy: 0.7999
Epoch 00003: val_accuracy improved from 0.84029 to 0.84694, saving model to checkpoints_best_only/checkpoint

INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 188 285us/sample - loss: 0.6721 - accuracy: 0.8001 - v

al_loss: 0.5140 - val_accuracy: 0.8469
Epoch 4/30
62200/62268 [=====] - ETA: 0s - loss: 0.6138 - accuracy: 0.8183
Epoch 00004: val_accuracy improved from 0.84694 to 0.86559, saving model to checkpoints_best_only/checkpoint

INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 188 285us/sample - loss: 0.6134 - accuracy: 0.8185 - v

al_loss: 0.4450 - val_accuracy: 0.8656
Epoch 5/30
62200/62268 [=====] - ETA: 0s - loss: 0.5707 - accuracy: 0.8304
Epoch 00005: val_accuracy did not improve from 0.86559

62268/62268 [=====] - 168 265us/sample - loss: 0.5712 - accuracy: 0.8304 - v

al_loss: 0.5034 - val_accuracy: 0.8560
Epoch 6/30
62200/62268 [=====] - ETA: 0s - loss: 0.5482 - accuracy: 0.8373 - v

al_loss: 0.5120 - val_accuracy: 0.8422
Epoch 7/30
62200/62268 [=====] - ETA: 0s - loss: 0.5278 - accuracy: 0.8437
Epoch 00007: val_accuracy improved from 0.86559 to 0.87870, saving model to checkpoints_best_only/checkpoint

INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 188 291us/sample - loss: 0.5278 - accuracy: 0.8437 - v

al_loss: 0.4081 - val_accuracy: 0.8787
Epoch 8/30
62200/62268 [=====] - ETA: 0s - loss: 0.5132 - accuracy: 0.8501
Epoch 00008: val_accuracy did not improve from 0.87870

62268/62268 [=====] - 168 265us/sample - loss: 0.5128 - accuracy: 0.8503 - v

al_loss: 0.5222 - val_accuracy: 0.8438
Epoch 9/30
62200/62268 [=====] - ETA: 0s - loss: 0.4982 - accuracy: 0.8541
Epoch 00009: val_accuracy improved from 0.87870 to 0.89317, saving model to checkpoints_best_only/checkpoint

INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 198 303us/sample - loss: 0.4981 - accuracy: 0.8541 - v

al_loss: 0.3761 - val_accuracy: 0.8892
Epoch 10/30
62200/62268 [=====] - ETA: 0s - loss: 0.4845 - accuracy: 0.8583
Epoch 00010: val_accuracy improved from 0.88916 to 0.93171, saving model to checkpoints_best_only/checkpoint

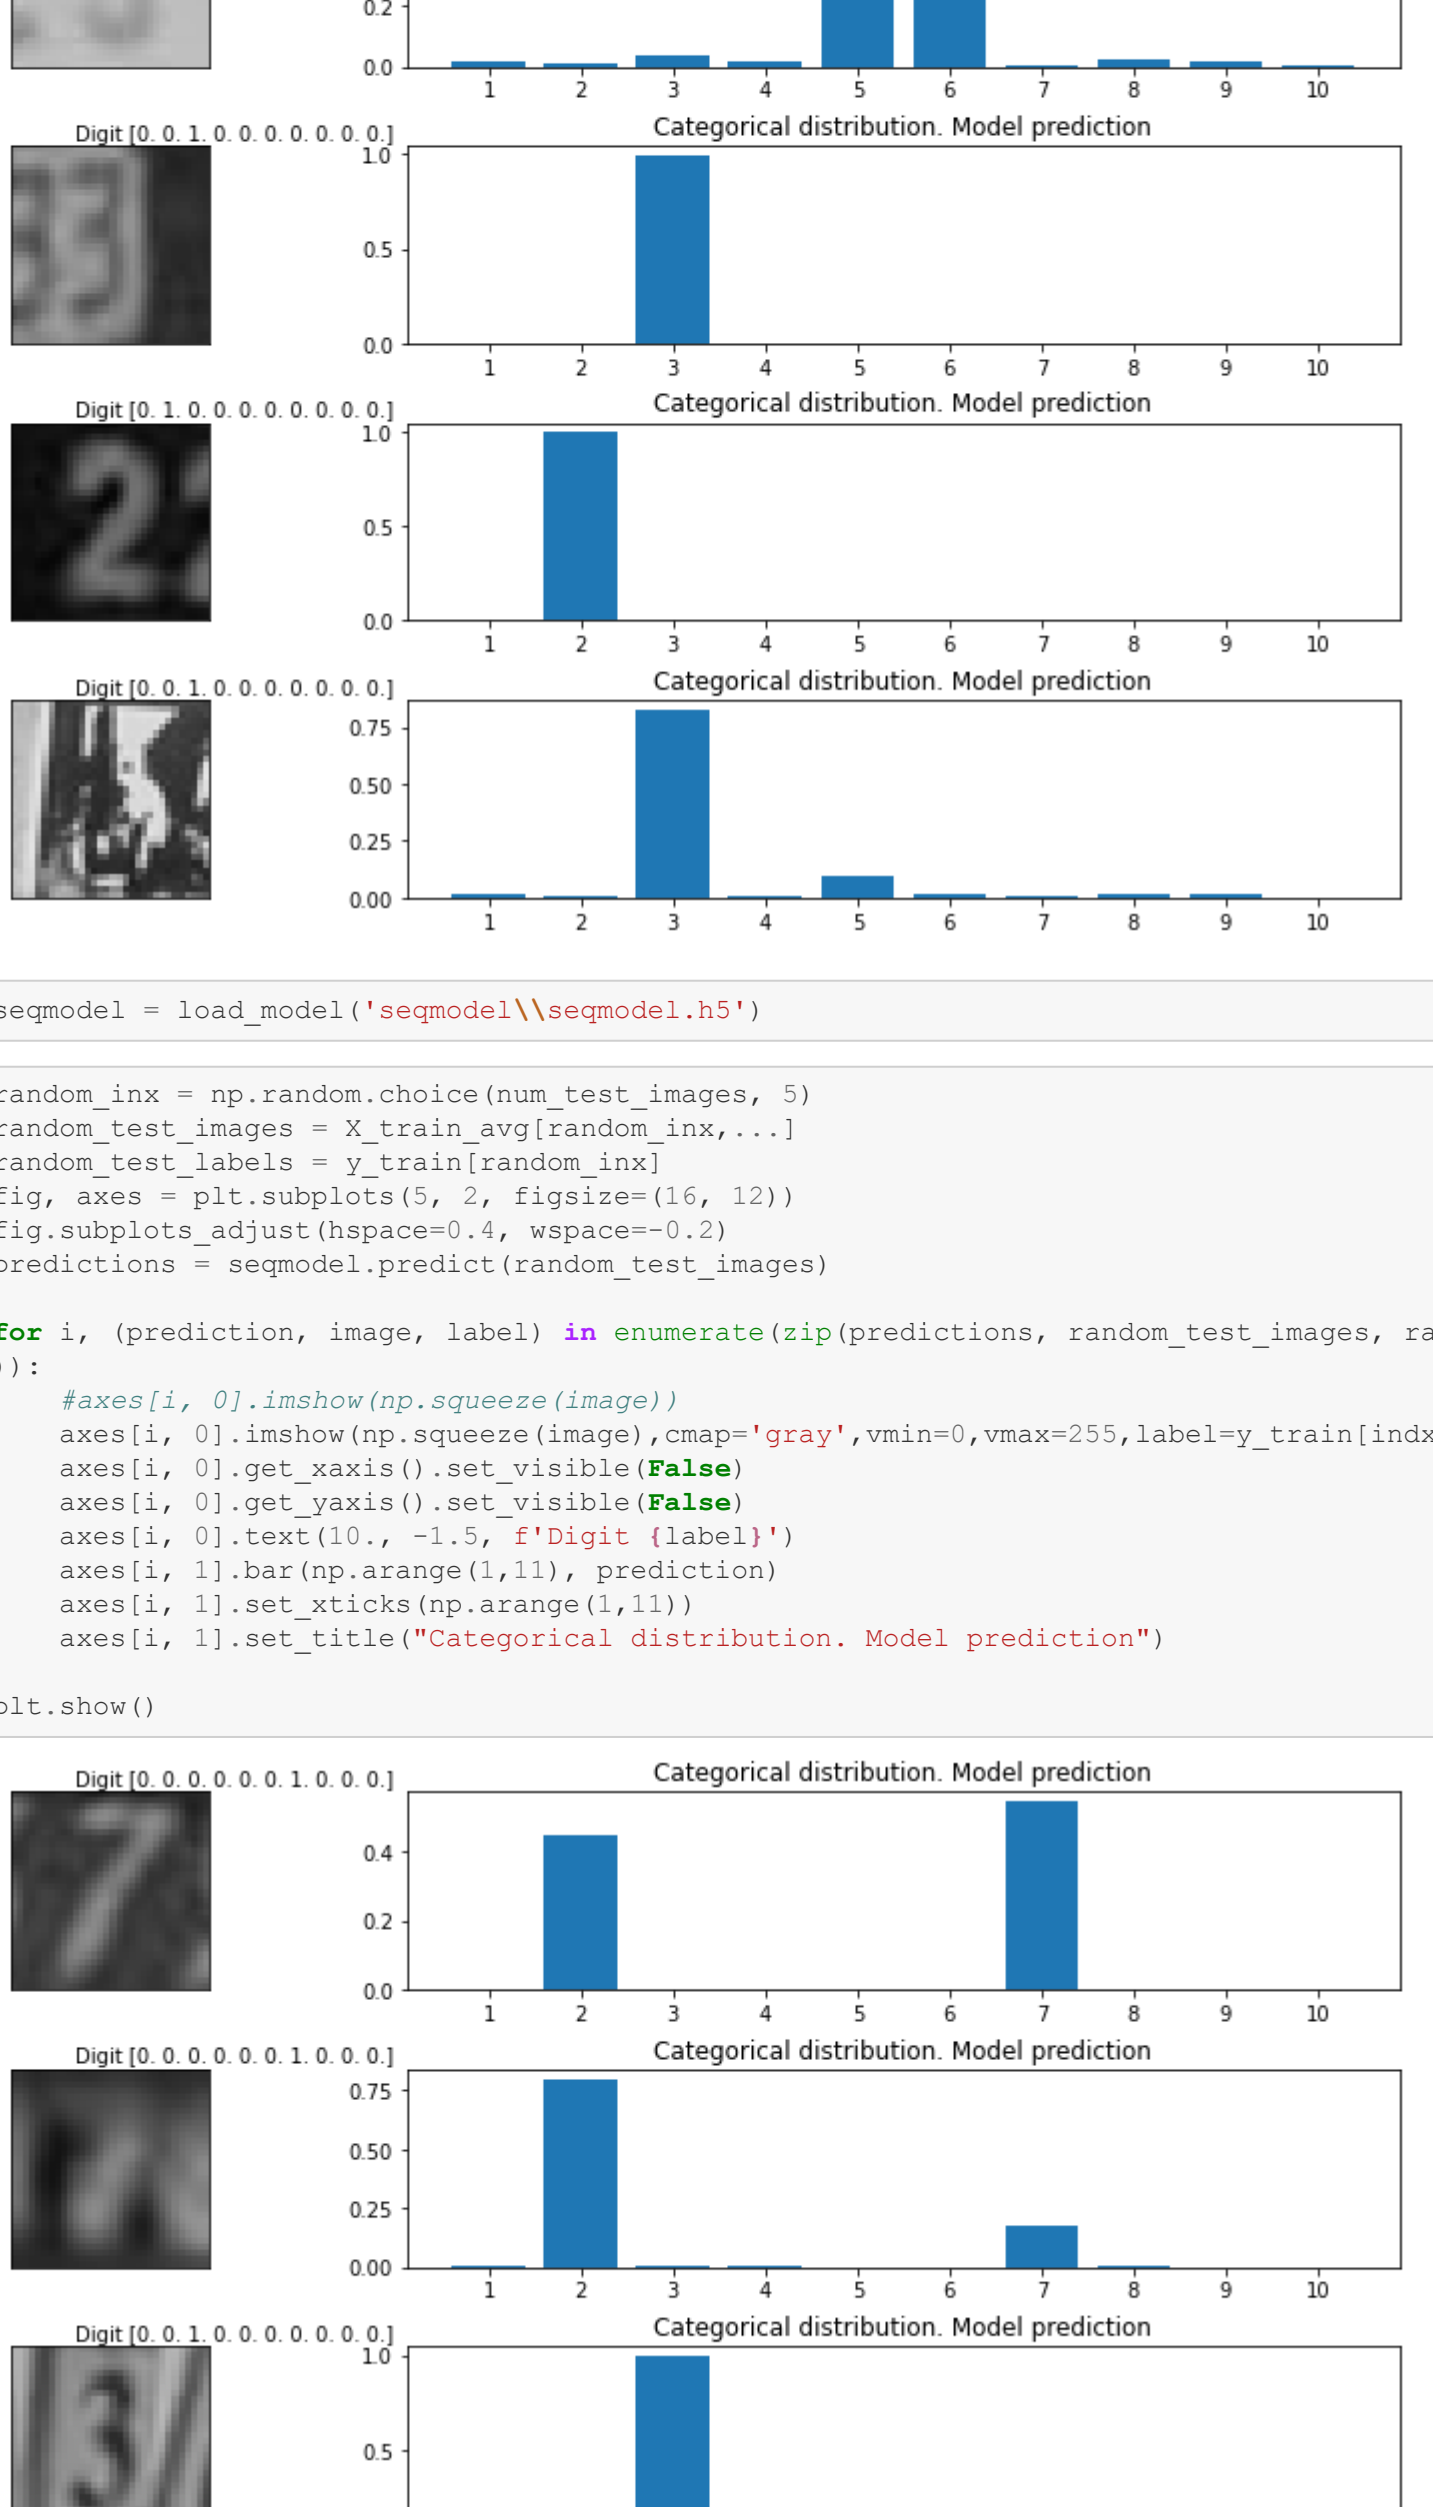
INFO:tensorflow:Assets written to: checkpoints_best_only/checkpoint/assets
62268/62268 [=====] - 218 345us/sample - loss: 0.4846 - accuracy: 0.8583 - v

al_loss: 0.3713 - val_accuracy: 0.8917
Epoch 11


```
[34]: random_inx = np.random.choice(num_test_images, 5)
random_test_images = X_train_avg[random_inx,...]
random_test_labels = y_train[random_inx]
fig, axes = plt.subplots(5, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=0.2)
predictions = cnnmodel.predict(random_test_images)

for i, (prediction, image, label) in enumerate(zip(predictions, random_test_images, random_test_labels)):
    #axes[i, 0].imshow(np.squeeze(image))
    axes[i, 0].imshow(np.squeeze(image), cmap='gray', vmin=0, vmax=255, label=y_train[indx])
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10, -1.5, f'Digit {label}')
    axes[i, 1].bar(np.arange(1,11), prediction)
    axes[i, 1].set_xticks(np.arange(1,11))
    axes[i, 1].set_title("Categorical distribution. Model prediction")

plt.show()
```



```
In [35]: seqmodel = load_model('seqmodel\\seqmodel.h5')
```

```
In [36]: random_inx = np.random.choice(num_test_images, 5)
random_test_images = X_train_avg[random_inx,...]
random_test_labels = y_train[random_inx]
fig, axes = plt.subplots(5, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=0.2)
predictions = seqmodel.predict(random_test_images)

for i, (prediction, image, label) in enumerate(zip(predictions, random_test_images, random_test_labels)):
    #axes[i, 0].imshow(np.squeeze(image))
    axes[i, 0].imshow(np.squeeze(image), cmap='gray', vmin=0, vmax=255, label=y_train[indx])
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10, -1.5, f'Digit {label}')
    axes[i, 1].bar(np.arange(1,11), prediction)
    axes[i, 1].set_xticks(np.arange(1,11))
    axes[i, 1].set_title("Categorical distribution. Model prediction")

plt.show()
```



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```